

DAY 8: LINEAR REGRESSION (AGENTIC AI CHALLENGE)

WHAT IS LINEAR REGRESSION?

Linear Regression is a supervised machine learning algorithm used to predict a continuous outcome based on one or more input features. It works by fitting a straight line (or curve, in polynomial regression) to data points that best explains the relationship between the independent and dependent variables.

WHY USE LINEAR REGRESSION?

Simple and interpretable.

Widely used for forecasting and trend analysis.

Acts as the foundation for more advanced algorithms.

WHERE IS IT USED?

Predicting house prices.

Estimating salary based on experience.

Forecasting demand, stock prices, etc.

CORE CONCEPT: BEST-FIT LINE USING LEAST SQUARES

We want to find the line $y = wX + b$ that minimizes the squared error:

FORMULA:

Loss Function (MSE):

$$J(w, b) = \frac{1}{n} \sum_{i=1}^n (y_i - (w x_i + b))^2$$

GOAL:

Minimize this function to find optimal w (weight/slope) and b (bias/intercept).

TYPES OF LINEAR REGRESSION

SIMPLE LINEAR REGRESSION

1 Feature (X), 1 Target (Y)

Fits a straight line

```
model.predict([[5]]) # Predict salary for 5 years experience
```

MULTIPLE LINEAR REGRESSION

Multiple features (X1, X2, ..., Xn)

Equation: $y = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b$

```
model.predict([[3000, 3, 2]]) # area, bedrooms, stories
```

POLYNOMIAL REGRESSION

Captures non-linear relationships by adding powers of X.

Still linear in terms of coefficients.

```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=2)
```

```
X_poly = poly.fit_transform(X)
model.fit(X_poly, y)
```

4 REGULARIZED LINEAR REGRESSION

Used to avoid overfitting by penalizing large coefficients.

◆ RIDGE REGRESSION (L2):

$$J(w) = \text{MSE} + \lambda \sum w_i^2 \quad J(w) = \text{MSE} + \lambda \sum w_i^2$$

◆ LASSO REGRESSION (L1):

$$J(w) = \text{MSE} + \lambda \sum |w_i| \quad J(w) = \text{MSE} + \lambda \sum |w_i|$$

```
from sklearn.linear_model import Ridge, Lasso
model = Ridge(alpha=1.0) # or Lasso(alpha=1.0)
```



EVALUATION METRICS

R^2 Score: How well the model explains variance in data.

Mean Squared Error (MSE): Average squared difference between actual and predicted values.



OVERFITTING VS UNDERFITTING



OVERFITTING

Model learns noise and patterns in training data.

High accuracy on training data, poor on test data.

Happens when the model is too complex (e.g., high-degree polynomial).



EXAMPLE:

Polynomial regression with degree = 10 on small data.

R^2 score is high on train but low on test.

▼ UNDERFITTING

Model fails to learn the pattern.

Poor performance on both training and test data.

Happens when the model is too simple.

🔍 EXAMPLE:

Linear regression on a curved dataset.

Predictions are not close to the actual values.

✅ IDEAL MODEL:

Balanced bias and variance.

Generalizes well on unseen data.

✅ SUMMARY TABLE

Type	Use-case	Code	Key Metric
Simple	1 feature	<code>LinearRegression()</code>	R^2 , MSE
Multiple	Multiple features	<code>LinearRegression()</code>	R^2 , MSE
Polynomial	Non-linear data	<code>PolynomialFeatures()</code>	R^2
Regularized	Avoid overfitting	<code>Ridge</code> , <code>Lasso</code>	R^2 , MSE

📁 DATASET EXAMPLES

Simple Regression: [Salary Dataset](#)

Multiple Regression: [Housing Dataset](#)