

DAY 10: OPTIMIZATION, GRADIENT DESCENT, MSE, R^2 & MULTIPLE LINEAR REGRESSION

GOAL OF THE DAY

Understand how optimization works using Gradient Descent, how it's powered by Calculus, evaluated using Mean Squared Error (MSE) and R^2 Score, and extended to Multiple Linear Regression for more complex predictions.

WHAT IS GRADIENT DESCENT?

Gradient Descent is an optimization algorithm used to find the minimum of a function — in ML, it minimizes the cost function.

WHY USE IT?

In linear regression, we want to minimize the cost:

$$\text{Cost (MSE)} = (1/n) * \sum (y - \hat{y})^2$$

But we don't solve it manually when data is large — instead, Gradient Descent updates weights iteratively:

UPDATE RULE:

$$\begin{aligned} w &= w - \alpha * \partial L / \partial w \\ b &= b - \alpha * \partial L / \partial b \end{aligned}$$

Where:

α = learning rate

$\partial L / \partial w$ = partial derivative of loss w.r.t. weight

STEPS OF GRADIENT DESCENT:

Initialize weights w and b randomly

Calculate predictions: $\hat{y} = wx + b$

Compute cost (MSE)

Compute gradients

Update w and b

Repeat until convergence

MSE AND R^2 IN DEPTH



MEAN SQUARED ERROR (MSE)

Measures average squared difference between actual and predicted values:

$$\text{MSE} = (1/n) * \sum (y - \hat{y})^2$$

Lower MSE \rightarrow better model

Used as a cost function in regression

R-SQUARED (R^2 SCORE)

Represents how much variance in target Y is explained by input X :

$$R^2 = 1 - (SS_{\text{res}} / SS_{\text{tot}})$$

Where:

$$SS_{\text{res}} = \sum (y - \hat{y})^2$$

$$SS_{\text{tot}} = \sum (y - \text{mean}(y))^2$$

✓ R^2 close to 1: Good fit

✓ R^2 close to 0: Bad fit



SIMPLE LINEAR REGRESSION EXAMPLE

```
from sklearn.linear_model import LinearRegression
import numpy as np
import pandas as pd

# Dataset
data = pd.read_csv("Salary_dataset.csv")
X = np.array(data['YearsExperience']).reshape(-1, 1)
y = data['Salary']

# Model
model = LinearRegression()
model.fit(X, y)

# Predict
print("Predicted salary for 5 years experience:",
      model.predict([[5]]))
```



OUTPUT

Coefficient (w): model.coef_

Intercept (b): model.intercept_

R-squared: model.score(X, y)

MULTIPLE LINEAR REGRESSION EXAMPLE

```
# Dataset: Housing.csv with ['area', 'bedrooms',  
'stories']  
X = np.array(data[['area', 'bedrooms', 'stories']])  
y = data['price']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2)  
model = LinearRegression()  
model.fit(X_train, y_train)  
  
print("R^2 score:", r2_score(y_test,  
model.predict(X_test)))
```

INTERPRETATION

Predict house price using multiple features

Visualize actual vs predicted prices

SUMMARY TABLE

Concept	Description
Gradient Descent	Optimizes parameters to minimize cost
MSE	Average squared error between actual and predicted
Learning Rate (α)	Controls step size during updates
Derivatives	Calculus tool to find slope for parameter updates
R ² Score	Measures model accuracy (0 to 1 scale)
Multiple Regression	Extends linear regression to multiple input variables



FINAL TAKEAWAY

“Gradient Descent is the heart of machine learning optimization — powered by calculus, evaluated by MSE, and validated by R^2 — it helps us build accurate and reliable prediction models.”